

# Data Collection from Web Sources

Session 2 @ CBA Batch 12

April 2019

sudhir\_voleti@isb.edu



1

## Background to the Example: Evolv-Xerox

- When looking for workers to staff its call centers, Xerox Corp. used **traditional hiring methods** and valued past experience on the job.
- Then, a computer program told Xerox Inc. that **experience doesn't matter**.
- Evolv Inc., a San Francisco start-up designed this algorithm that crunches and matches *psychographic* profiles with work outcomes.
- Claim was that conventional Hiring methods are remarkably short on rigor for this crucial business function
- Can a **statistical approach to hiring** improve results by reducing hiring managers' biases?



2

### Background to the Example: Evolv-Xerox

- Evolv's Method: Job applicants take a 30-minute test that screens them for personality traits.
- Qs include Agree-Disagree on a 1-7 scale for "I ask more questions than most people do.", "People tend to trust what I say." Etc.
- The firm has a **database** of **current and past** worker psychographics and job performance.
- It '*statistically*' *matches* the psychographic patterns and looks for the best job performance prediction for each application.
- It then assigns a score - Green, Yellow or Red - to each applicant.



3

### Evolv-Xerox Example: So what happened finally?

- Attrition dropped 20%. On-job performance improved similarly.
- The algo uncovers correlations the naked eye can't see. Some fun facts the algo uncovered...
- "Who would you rather hire – A or B?"
- A) Candidate who used Internet Explorer or Safari to answer your questionnaire
- B) Candidate who used Firefox or Chrome?
- A) Candidate who has joined 1-2 social networks
- B) Candidate who has joined 3-4 social networks?



4

## Evolv-Xerox: Inputs to the machine

- So, how did a set of humble Likerts provide inputs for this fancy analysis?
- Consider the data Evolv had: **X**
  - the individual side (demographics, psychographics, etc.)
  - the firm side (training costs, productivity & turnover rates etc.) **Y**
- Given this data, Evolv is somehow able to *connect* job outcomes to psychographic profiles... and thereafter *predict* outcomes for new psychographic profiles... **f(.)**
- So, what *principle* drives Evolv's approach?

$$y = f(X) + e$$



5

## Motivating Example: Concluded

- There's a feeling that survey research is so 20<sup>th</sup> century.
- Pen-and-paper. Grunt work. Low-tech.
- That may be so.
- However, the primary DC it provides can yield insight-gold.
- Even humble Likert Qs on psyche-related phenomena can sometimes yield very powerful predictors (as latent feature combinations) ...
- ... that today's powerful machine learning algos can happily munch on.



6

## Session Outline

- A Motivating Example
  - Evolv-Xerox and Likert Scales
- Basic HTML primer
- DC from static web pages
  - Examples 1, 2, 3
- DC from dynamic webpages
  - Using webdrivers in py
- Session Wrap-up

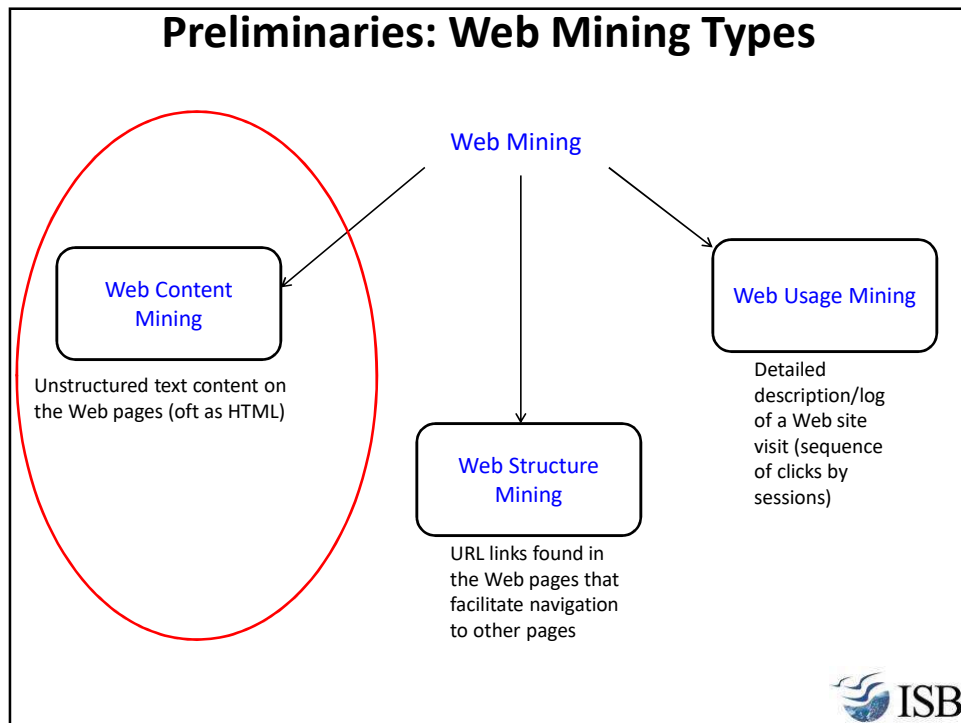


7

## Some Preliminaries



8



9

### Preliminaries: HTML & the DOM

- What is *web-scraping*? What is its objective?
- What are webpages written in?
- What is HTML? What does it look like?
- What is a DOM? What does it look like?
- What converts HTML & CSS to human-readable form?
- How much of HTML code would typically be of interest in web-scraping?

ISB

10

### Preliminaries: Static and Dynamic web pages

- Part of a web-page may have merely 'static' content in html.
- Part of the page may be **dynamic** - loading and updating constantly. (E.g., comments sections, twitter or FB timelines etc)
- Which of the 2 would be easier to scrape?
- For static pages, we'll use *rvest* in R and *beautiful soup* in Py.
- Sometimes we may need to evaluate conditions interactively and then only decide what to scrape and when. Enter *webdrivers*.
- Before we proceed, a basic primer on HTML ...



11

## A Basic HTML Primer



12

### **An Elementary HTML primer**

- Aim of this primer is to demonstrate:
- Basic html structures
- Major html tag-types such as:
  - a. title tag
  - b. body tags
  - c. document structure tags
  - d. basic markup tags
  - e. links, image tags
  - f. table tags
- Idea is to enable easy scanning of the [page-source](#) of static webpages.
- Open file '[html primer dc text-editor.txt](#)' and *save-as* a '.html' file.  
Now open the .html file in browser.



13

### **Basic HTML Primer: Recap**

- List some HTML tags we just saw.
- Consider the power and possibility to create variously designed and marked-up webpages simply remixing a handful of basic HTML tags.
- Can we now see page-sources on the web and estimate what the actual markups may look like?
- Ready to handle simple homework on building a static web-page?
- For scraping data off of web-pages, we'll need to identify particular patterns in html tags.
- Let's head there next, then.



14

## Scraping Static Web-pages

rvest and BeautifulSoup



15

### A Roadmap for this Section

- Will need CSS selector gadget on chrome browser for this section.
- Plan is to cover 3 examples.
- 1. Scraping the IMDB - mini use-case as a walkthrough example.
- 2. Scraping a Telecom glossary @ [www.atis.org/glossary](http://www.atis.org/glossary)
- 3. Py scraping news-stories @ [news.ycombinator.com](http://news.ycombinator.com)
- Ready? Open '[rvest\\_example.html](#)'



16



## A Walk-through Example: IMDB

- Q: “What attributes characterize the *best* movies of all time?”
- How would you approach this Q? Where would you first look for data?
- Even before DC, what might “attributes” here mean? And what is ‘best’ anyway?



17

## rvest for web scraping: Walkthrough Example

- Step 1: Go to the webpage, study it. ID features of interest such as

IMDb Charts

### Top Rated Movies

Top 250 as rated by IMDb Users

Showing 250 Titles

Sort by: Ranking

Rank & Title	IMDb Rating	Your Rating
1. The Shawshank Redemption (1994)	★ 9.2	+
2. The Godfather (1972)	★ 9.2	+
3. The Godfather: Part II (1974)	★ 9.0	+
4. The Dark Knight (2008)	★ 8.9	+
5. 12 Angry Men (1957)	★ 8.9	+



18

## rvest for web scraping: Walkthrough Example

- Step 2: Read it into R using \_\_\_\_\_ function.

```
library("rvest")
library("XML")
# IMDB Top 250 Movies
url = "http://www.imdb.com/chart/top?ref_=nv_wl_img_3"
page = read_html(url)
```

- Step 3: Activate the SelectorGadget tool on Chrome browser.




19

## rvest for web scraping: Walkthrough Example

- Step 4: ID the CSS tags needed and read them into \_\_\_\_\_ function.

```
movie.nodes = html_nodes(page, '.titleColumn a')
# Check one node
xmlTreeParse(movie.nodes[[1]])
```



20

- Step 5: Scan the structure of extracted nodes. ID that tiny % that is useful using the \_\_\_\_\_ function.

```
## $doc
## $file
## [1] "<buffer>"
##
## $version
## [1] "1.0"
##
## $children
## $children$a
## <a href="/title/tt0111161/?pf_rd_m=A2FGELUUNOQUNL&pf_rd_p=2398042102&pf_rd_r=01F87A7YM2837BHDB905&pf_rd_s=center-1&pf_rd_t=1506&pf_rd_i=top&ref_=chttp_tt_1" title="Frank Darabont (dir.), Tim Robbins, Morgan Freeman">The Shawshank Redemption</a>
##
##
## attr(,"class")
## [1] "XMLDocumentContent"
##
```

B

- Step 6: Store fields of interest as separate vectors because \_\_\_\_.

```
> head(cbind(movie.link, movie.cast, movie.name))
  movie.link
[1] "http://www.imdb.com/title/tt0111161/?pf_rd_m=A2FGELUUNOQJNL&pf_rd_p=2398042102&pf_rd_i=top&ref=chttp_tt_1"
[2] "http://www.imdb.com/title/tt0068646/?pf_rd_m=A2FGELUUNOQJNL&pf_rd_p=2398042102&pf_rd_i=top&ref=chttp_tt_2"
[3] "http://www.imdb.com/title/tt0071562/?pf_rd_m=A2FGELUUNOQJNL&pf_rd_p=2398042102&pf_rd_i=top&ref=chttp_tt_3"
[4] "http://www.imdb.com/title/tt0468569/?pf_rd_m=A2FGELUUNOQJNL&pf_rd_p=2398042102&pf_rd_i=top&ref=chttp_tt_4"
[5] "http://www.imdb.com/title/tt0050083/?pf_rd_m=A2FGELUUNOQJNL&pf_rd_p=2398042102&pf_rd_i=top&ref=chttp_tt_5"
[6] "http://www.imdb.com/title/tt0108052/?pf_rd_m=A2FGELUUNOQJNL&pf_rd_p=2398042102&pf_rd_i=top&ref=chttp_tt_6"
  movie.cast      movie.name
[1] "Frank Darabont (dir.), Tim Robbins, Morgan Freeman" "The Shawshank Redemption"
[2] "Francis Ford Coppola (dir.), Marlon Brando, Al Pacino" "The Godfather"
[3] "Francis Ford Coppola (dir.), Al Pacino, Robert De Niro" "The Godfather: Part II"
[4] "Christopher Nolan (dir.), Christian Bale, Heath Ledger" "The Dark Knight"
[5] "Sidney Lumet (dir.), Henry Fonda, Lee J. Cobb" "12 Angry Men"
[6] "Steven Spielberg (dir.), Liam Neeson, Ralph Fiennes" "Schindler's List"
```

11

## rvest for web scraping: Walkthrough Example

- Step 7: Sometimes some basic \_\_\_\_\_ maybe needed to clean up the info before storing. Take movie release year ....

```
year = gsub(")", "",
           gsub("\\(", "",
               html_text(
                 html_nodes(page, '.secondaryInfo')
               )))
```

.secondaryInfo

Clear (250)

Toggle Position



23

## rvest for web scraping: Walkthrough Example

- Step 7: Sometimes some basic \_\_\_\_\_ maybe needed to clean up the info before storing. Take **average movie rating** that has CSS tags \_\_\_\_\_.

```
rating.nodes = html_nodes(page, '.imdbRating')
# Check One node
xmlTreeParse(rating.nodes[[1]])
```

```
## $children
## $children$td
## <td class="ratingColumn imdbRating">
## <strong title="9.2 based on 1,707,259 user ratings">9.2</strong>
## </td>
##
```

```
# Correct the node
rating.nodes = html_nodes(page, '.imdbRating strong')
votes = as.numeric(gsub(' ', '',
                       gsub(' user ratings', '',
                           gsub('.*?based on ', '',
                               sapply(html_attrs(rating.nodes), `[`, 'title')
                             )))
rating = as.numeric(html_text(rating.nodes))
```



24

## rvest for web scraping: Walkthrough Example

- Step 8: Build a dataframe with all extracted info and write to disk.

```
top250 = data.frame(movie.name, movie.cast, movie.link, year, votes, rating)
head(top250)
```

##	movie.name	##	year	votes	rating
## 1	The Shawshank Redemption	## 1	1994	1707259	9.2
## 2	The Godfather	## 2	1972	1167246	9.2
## 3	The Godfather: Part II	## 3	1974	799349	9.0
## 4	The Dark Knight	## 4	2008	1694374	8.9
## 5	Schindler's List	## 5	1993	874184	8.9
## 6	12 Angry Men	## 6	1957	453555	8.9

```
##
##      movie.cast
## 1  Frank Darabont (dir.), Tim Robbins, Morgan Freeman
## 2  Francis Ford Coppola (dir.), Marlon Brando, Al Pacino
## 3  Francis Ford Coppola (dir.), Al Pacino, Robert De Niro
## 4  Christopher Nolan (dir.), Christian Bale, Heath Ledger
## 5  Steven Spielberg (dir.), Liam Neeson, Ralph Fiennes
## 6  Sidney Lumet (dir.), Henry Fonda, Lee J. Cobb
##
movie.link
## 1 http://www.imdb.com/title/tt0111161/?pf_rd_m=A2FGELUUNOQJN
5&pf_rd_s=center-1&pf_rd_t=15506&pf_rd_i=top&ref_=chttp_tt_1

write.csv(top250, 'IMDB Top 250.csv', row.names = F)
```



25

## rvest for web-scraping: Recap

- We saw just how easy it can be to scrape data off (reasonably structured) web-pages.
- But hey, suppose you really needed to collect **movie genre**, then what?
- Turns out 'movie.link' links to the movie's IMDB page where such info is available.
- One can replicate the previous analysis this time on the movie's IMDB page --> ID & extract more features --> loop over the 250 movies.
- Heck, if there are more links of interest, then one can "chain" together web scraping ops on successive pages.



26

### rvest : Recap and Ponder

- A good time to take a step back and review learnings from rvest
- What libraries did we call?
- What main inbuilt functions did we use?
- What user-defined functions did we use?
- What inputs and outputs to the web-scraping did we see?
- Any other comment on learnings? Applications? Assignments?



27

### Another rvest Exercise

- So far we built vectors of **features of interest** (e.g., movie name, cast etc) by *applying* funcs **over the entire webpage**, then *column-binding* into DFs.
- Alternately, we could've taken each unit (e.g., movie) as an R object, **extracted features** of interest and *row-bound* the result into a DF.
- Open the HTML file '**Webscraping Basics with R and Py**'.
- The file contains 2 examples - one for scraping a telecom industry glossary page-by-page, and the other for scraping the ycombinator main news page.



28

## Scraping a Glossary: Example

- Why scrape atis.org?
  - Well, this exercise is illustrative. For the tools & skills that're more generally applicable.
- Note [change in URLs](#) as we click on each link.
- In [structured](#) web-pages, there is a **pattern** to that URL change that we can leverage to scrape the site easily.
- Each glossary term in turn is a [link](#), with an [id](#) that leads to a new page (check its URL) with the text definition of that term.
- Our aim scrape the glossary's terms alone.
  - But we could, if need be, scrape the text definition as well.



29

## Scraping a Glossary: Recap

- What did we just do?
- What libraries did we call?
- What functions did we use?
- What kind of structure did we rely on?
- P.S. What if we also needed to capture each glossary term's text definition?



30

### Scraping news.ycombinator: Another Exercise

- Along the same lines, however there's a twist here.
- This one requires we study the view-source more deeply.
- From page-source, we can id the nodes of interest (e.g., 'a') and the attributes in that node (e.g., 'storylink', 'sitestr') that contain quantities of interest like text, links etc.
- Let's do this one in Python's [Beautiful Soup](#).
- Open [spyder](#) and let's begin.



31

### Beautiful soup-ing the webpage in py

- Py up and running for everyone? Open spyder (or Jupyter).
- Copy paste the code, chunk by chunk.
- Examine the variables get created in the variable explorer pane.
- Recall that the hard work of navigating the page-source was done already.



32



## A Few More Preliminaries

Data Display vs Data Storage formats



33

### Some Well-known Data Storage Formats

- Unlike HTML & DOM which are more of data *markup and display* formats, JSON and XML are popular data *storage* formats.
- Consider an example of fields {Name, age, occupation} for two people A & B.
  - {Ravi, 38, Graphic Designer}
  - {Anu, Sales Executive, 27}
- Consider how a person vs how a machine would read & understand.
  - Why the difference? What can be done about *ensuring* such doesn't happen?
- Enter data storage formats like JSON and XML.
  - These contain both the field names and the field values for every data point.
  - Verbose, but accurate.
  - Sample this example



34

## JSON & XML Data Storage Formats

- Here's a quick view of what JSON output looks like ...
- Can you ID the field names (or 'keys') and values?
- And now a quick view of what XML output looks like ...
- Can you ID the field names (or 'keys') and values?
- Note the ability to nest and build hierarchical data storage structures



35

## Webdrivers for Web-scraping



36

## Using Selenium in Py

- Open the HTML file '[Intro to webdrivers - Selenium in Py](#)'
- Let's walk through it step by step.
- What I'll show next is fairly basic. However, ...
- If you're aware of alternatives, better ways to do the same thing etc., pls speak up and share with the class.



37

## Using Webdrivers: Recap

- What are webdrivers and where are they most used?
- What modules did we invoke for using webdrivers?
- What main functions were called? What did they do?
- What further possibilities come to mind with webdriver use?
- Ready for some basic homework involving py and selenium?



38

## Session Wrap-up



39

### Session Wrap-up

- This was more of a 'workshop' session than the 2 preceding it.
- Expect the trend to continue here onwards for the rest of DC and for Text-analytics.
- Group formation issues remain.
- Any Qs or comments etc?



40