

Enzyme Classifier - SparkML

July 26, 2019

Team Members

- Shubhendu Vimal - 11915067
- Dharani Kiran Kavuri - 11915033
- Anmol More - 11915043

ReadMe :

- Data Preparation through python script
- Convert raw data csv and then to libsvm in R
- Run spark ML algos in jupyter notebook

```
In [1]: import pandas as pd
import glob
from pyspark.ml.classification import LogisticRegression
from pyspark.ml import Pipeline
from pyspark.ml.classification import DecisionTreeClassifier
from pyspark.ml.classification import RandomForestClassifier
from pyspark.ml.classification import MultilayerPerceptronClassifier, OneVsRest
from pyspark.ml.feature import IndexToString, StringIndexer, VectorIndexer
from pyspark.ml.evaluation import MulticlassClassificationEvaluator
from pyspark.sql import SparkSession
```

Read data

Read cleaned and processed data for enzymes from data directory

```
In [2]: data = pd.DataFrame()
files = glob.glob("data/*.csv")
for file in files :
    print("Processing File :", file)
    df = pd.read_csv(file)
    print(df.shape)
    data = data.append(df, ignore_index=True)

def type_to_numeric(x):
    if x=='uniprot-Isomerase':
```

```

        return 0
    if x=='uniprot-Lyase':
        return 1
    if x=='uniprot-ligase':
        return 2
    if x=='uniprot-Transferases' :
        return 3
    if x=='uniprot-Hydrolases' :
        return 4
    if x=='uniprot-Oxidoreductases' :
        return 5

data['type_numeric'] = data['Type'].apply(type_to_numeric)

#Final csv to work upon and run spark ML algorithms
#Taking 100000 records of each enzyme
data.to_csv("final_data/protien-sequences.csv", index=False)
data.sample(5)

```

```

Processing File : data/uniprot-Oxidoreductases.csv
(89604, 24)
Processing File : data/uniprot-Isomerase.csv
(100001, 24)
Processing File : data/uniprot-Hydrolases.csv
(100001, 24)
Processing File : data/uniprot-ligase.csv
(100001, 24)
Processing File : data/uniprot-Lyase.csv
(100001, 24)
Processing File : data/uniprot-Transferases.csv
(100001, 24)

```

```

Out[2]:
      A      C      D      E      F      G      H      I      K      L      ...      S      \
334374  2.0  0.0   5.0  12.0   2.0   7.0   1.0  10.0   7.0  10.0   ...   7.0
447198 22.0   1.0   7.0   7.0   6.0  14.0   8.0  10.0   0.0  14.0   ...   6.0
329103 18.0   3.0  12.0  15.0  11.0  21.0   5.0  19.0  16.0  33.0   ...  24.0
215190 31.0   7.0  11.0  19.0  10.0  14.0   8.0  11.0   5.0  35.0   ...  14.0
92571  22.0   4.0  19.0  20.0  10.0  20.0   8.0  15.0  22.0  25.0   ...  17.0

```

```

                                Sequence      T      \
334374  MGRTDDMLIIRGVNVFPSQIESVLLENGDTTPHYQLIVNRKGNLDD...   5.0
447198  MPHRILVLHGPNLNLGTTREPEVYGRITTLADIDAALTAQAQTAGAE...   9.0
329103  MVHELLQAKWRIIDQSHFGPMFDAKQSFAIDDTLCTSVGKGLSDP...  13.0
215190  MFDIGVNLTSTQFAKDRDKVVKRAREAGISGMLITGTNALESQQAL...  14.0
92571   MKLSFNTWVYNSFPSMLPFYPLEEVISRIAAFGYDGIIGCASP...   9.0

```

```

Type      V      W      Y      gravity      weight      type_numeric

```

334374	uniprot-ligase	10.0	0.0	1.0	-0.410577	11844.5304	2
447198	uniprot-Lyase	11.0	1.0	5.0	0.094595	15959.8291	1
329103	uniprot-ligase	21.0	4.0	13.0	-0.012950	31174.5785	2
215190	uiprot-Hydrolases	13.0	5.0	2.0	-0.122394	28765.5914	4
92571	uniprot-Isomerase	15.0	8.0	14.0	-0.422857	31818.8477	0

[5 rows x 25 columns]

R Code

Convert csv data to libsvm for sparkML

In [3]: `%load_ext rpy2.ipython`

In [4]: `%%R`

```
# Ref : https://vatsalbits.wordpress.com/2016/01/13/csv-to-libsvm/

# download e1071 library if not available
if (!require(e1071)) {install.packages("e1071")}

# download sparseM library if not available
if (!require(SparseM)) {install.packages("SparseM")}

# load the libraries
library(e1071)
library(SparseM)

# load the csv dataset into memory
data <- read.csv('final_data/protien-sequences.csv')

# take the numeric columns are format as matrix
x <- as.matrix(data[,c('A', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'K', 'L', 'M',
                       'N', 'P', 'Q', 'R', 'S', 'T', 'V', 'W', 'Y',
                       'gravity', 'weight')])

# assign labels to vector y
y <- data[, 'type_numeric']

# convert input columns to sparse matrix
x_matrix <- as.matrix.csr(x)

# write output to libsvm format
write.matrix.csr(x_matrix, y=y, file="final_data/protien-sequences-libsvm.txt")
```

R[write to console]: Loading required package: e1071

R[write to console]: Loading required package: SparseM

R[write to console]:

Attaching package: SparseM

R[write to console]: The following object is masked from package:base:

backsolve

In [5]: *#create a new spark session*

```
MAX_MEMORY = "5g"
spark = SparkSession.builder.appName('protien-classifier').
config("spark.executor.memory", MAX_MEMORY).
config("spark.driver.memory", MAX_MEMORY).getOrCreate()
```

In [6]: *# Load training data*

```
data = spark.read.format("libsvm").load("final_data/protien-sequences-libsvm.txt")
```

Run Spark-Mllib Algorithms

Logistic Regression Classifier

Code chunks are directly taken from official spark-ml documentation.

Ref - <https://spark.apache.org/docs/latest/ml-classification-regression.html>

In [7]: *%%time*

```
#split dataset in test and training sets
(train, test) = data.randomSplit([0.7, 0.3])

#use max iteration of 10, and few default parameters
logistic_reg = LogisticRegression(maxIter=20, regParam=0.005,
                                   elasticNetParam=0.005)

# Fit the model
logic_reg_model = logistic_reg.fit(train)

# coefficients and intercept for multi class logistic regression
print("Logistic Reg Coefficients: " + str(logic_reg_model.coefficientMatrix))
print("Logistic Reg Intercept: " + str(logic_reg_model.interceptVector))

# Print model summary
model_summary = logic_reg_model.summary

accuracy = str(model_summary.accuracy)
false_positive_rate = str(model_summary.weightedFalsePositiveRate)
true_positive_rate = str(model_summary.weightedTruePositiveRate)
f_measure = str(model_summary.weightedFMeasure())
precision = str(model_summary.weightedPrecision)
recall = str(model_summary.weightedRecall)

print("Model accuracy : " + accuracy +
      "\nFalse positive rate : " + false_positive_rate +
```

```

"\nTrue positive rate : " + true_positive_rate +
"\nF-measure : " + f_measure +
"\nPrecision : " + precision +
"\nRecall : " + recall)

```

```

Logistic Reg Coefficients: DenseMatrix([[ 5.45202991e-03, -8.41721306e-02,  1.47792590e-02,
      3.31893586e-03, -3.73312494e-04, -9.73336894e-04,
     -1.93253095e-02, -5.61266461e-03,  2.89474262e-02,
      3.58842586e-03,  6.87395544e-04, -1.34753590e-03,
     -2.70238535e-03,  1.24997026e-02,  3.13580444e-04,
     -6.59144958e-03,  1.79010417e-03, -1.75218553e-03,
     -1.02059722e-01, -1.99632503e-02, -8.86660848e-01,
      8.33571514e-07],
 [ 8.85022716e-03, -1.18602161e-02,  2.83949155e-03,
     -4.75651902e-03, -1.67761010e-02,  3.44992374e-03,
     -2.00237860e-02, -3.42373036e-03, -1.04233788e-02,
      1.00570350e-03,  9.45213109e-03,  2.87557424e-02,
     -0.00000000e+00,  4.38811521e-03, -3.17804872e-03,
      7.87092740e-03,  1.46965982e-02, -3.95256967e-03,
     -3.74618000e-02, -7.53803470e-03,  1.78657822e-01,
      1.01338335e-06],
 [-5.44382414e-04, -3.65995056e-02,  1.58249042e-02,
      2.90505556e-02,  1.97166963e-02, -2.58209170e-03,
     -2.72139077e-02, -3.92640394e-03,  1.39468558e-02,
      6.41660339e-03, -4.90823743e-03, -1.22397775e-02,
     -2.10386699e-03, -2.55006639e-03,  3.58658471e-04,
     -9.47427993e-03, -3.44740078e-03, -1.08598272e-03,
      2.32355489e-02,  1.23598495e-02, -1.95403376e-01,
     -1.24076148e-06],
 [-2.03297917e-02,  1.21864343e-02, -1.58565824e-02,
     -2.09937573e-02, -1.88554774e-02,  6.02781399e-03,
      2.04237107e-02,  8.34704605e-02,  2.49272055e-03,
     -3.20490486e-02, -6.43147253e-02,  7.89723454e-03,
     -7.28943716e-03, -9.92204577e-03,  2.17382198e-03,
     -7.46535750e-04, -1.25944695e-02,  2.74750585e-02,
      1.58983343e-02, -0.00000000e+00,  1.12359426e+00,
     -3.18344785e-07],
 [-2.45574514e-03,  2.79102499e-02,  1.62842341e-03,
     -4.20275874e-03,  3.90743498e-03, -2.62320748e-03,
      6.86609432e-02, -6.35138103e-03, -4.75216675e-03,
      1.00420738e-03,  1.78526063e-02, -2.75581708e-03,
      7.65353346e-04, -5.89824624e-03, -3.73230530e-04,
      1.84585536e-02,  1.37432514e-03, -7.31473093e-03,
     -4.08326154e-03,  1.26975951e-02, -5.71198888e-02,
      4.24723852e-07],
 [-7.61661945e-04,  1.22353827e-01, -1.39996583e-02,
     -9.55126011e-03,  8.01425588e-03, -3.50855993e-03,
      8.79739093e-03, -1.08732097e-02, -9.43516283e-03,

```

```

0.00000000e+00, 1.20510455e-03, -7.56656920e-03,
7.07526911e-03, 6.41561717e-03, -7.95823268e-04,
2.42710127e-04, -4.77668365e-03, -6.11424855e-03,
3.65983003e-02, 6.71549501e-04, -3.34873437e-01,
-6.77072976e-07]])
Logistic Reg Intercept: [0.17462809241234534,0.010481053067311924,-0.6559807909752295,0.535728
Model accuracy : 0.36205831255588106
False positive rate : 0.12731148458569602
True positive rate : 0.36205831255588106
F-measure : 0.3611527231912624
Precision : 0.3628963004678636
Recall : 0.36205831255588106
CPU times: user 21.3 ms, sys: 6.57 ms, total: 27.8 ms
Wall time: 12.6 s

```

Decision Tree Classifier

```

In [8]: %%time
#mark labels as index of libsvm data
label_index = StringIndexer(inputCol="label", outputCol="index_label").fit(data)

# Mark all input columns as features
feature_index = VectorIndexer(inputCol="features",
                              outputCol="index_features",
                              maxCategories=2).fit(data)

# Lets do 60:40 ratio for training and test data
(train, test) = data.randomSplit([0.6, 0.4])

# Make a decision tree model
decision_tree_classifier = DecisionTreeClassifier(labelCol="index_label",
                                                  featuresCol="index_features")

# Create a pipeline of decision tree model
pipeline = Pipeline(stages=[label_index, feature_index, decision_tree_classifier])

# train our decision tree using input columns
decision_tree_model = pipeline.fit(train)

model_summary = decision_tree_model.stages[2]
# summary only
print("Created decision tree classifier with : \n")
print(model_summary)

# Get predictions for test set
predictions_test = decision_tree_model.transform(test)

```

```
# Look at top 10 predictions, 9/10 predictions seems correct
predictions_test.select("prediction", "index_label", "features").show(10)

# calculate accuracy on test set
multi_class_classification_evaluator = MulticlassClassificationEvaluator(
    labelCol="index_label", predictionCol="prediction", metricName="accuracy")
accuracy_percentage = multi_class_classification_evaluator.evaluate(predictions_test)

print("Test set error = %f " % (1.0 - accuracy_percentage))
```

Created decision tree classifier with :

DecisionTreeClassificationModel (uid=DecisionTreeClassifier_62c948558ae8) of depth 5 with 51 nodes

```
+-----+-----+-----+
|prediction|index_label|          features|
+-----+-----+-----+
|      5.0|      5.0|(22,[0,1,2,3,4,5,...|
|      5.0|      5.0|(22,[0,1,2,3,4,5,...|
|      0.0|      5.0|(22,[0,1,2,3,4,5,...|
|      3.0|      5.0|(22,[0,1,2,3,4,5,...|
|      5.0|      5.0|(22,[0,1,2,3,4,5,...|
|      5.0|      5.0|(22,[0,1,2,3,4,5,...|
|      5.0|      5.0|(22,[0,1,2,3,4,5,...|
|      5.0|      5.0|(22,[0,1,2,3,4,5,...|
|      3.0|      5.0|(22,[0,1,2,3,4,5,...|
|      3.0|      5.0|(22,[0,1,2,3,4,5,...|
+-----+-----+-----+
```

only showing top 10 rows

Test set error = 0.651715

CPU times: user 41.3 ms, sys: 9.75 ms, total: 51.1 ms

Wall time: 13.2 s

Random Forest Classifier

As per our evaluation, this gives us best result with accuracy of 57% with 30 trees

We can further work on this to increase number of trees, max depth and other hyper parameter tunings

```
In [9]: %%time
#mark labels as index of libsvm data
label_index = StringIndexer(inputCol="label", outputCol="index_label").fit(data)

# Mark all input columns as features
feature_index = VectorIndexer(inputCol="features", outputCol="index_features",
                              maxCategories=2).fit(data)

# Lets do 70:30 ratio for training and test data this time
```

```

(train, test) = data.randomSplit([0.7, 0.3])

# Create a random forest classifier with 20 trees
random_forest_classifier = RandomForestClassifier(labelCol="index_label",
                                                featuresCol="index_features",
                                                numTrees=30, maxDepth=10)

# label converter to convert indexed labels
label_converter = IndexToString(inputCol="prediction", outputCol="predicted_label",
                               labels=label_index.labels)

# Create a pipeline of random forest model
pipeline = Pipeline(stages=[label_index, feature_index,
                             random_forest_classifier, label_converter])

# Train random forest classifier
random_forest_model = pipeline.fit(train)

model_summary = random_forest_model.stages[2]
# summary only
print("Created random forest classifier with : \n")
print(model_summary)

# Get predictions for test set
predictions_test = random_forest_model.transform(test)

# see few of our predictions
predictions_test.select("predicted_label", "label", "features").show(10)

# Calculate test set error
multi_class_classification_evaluator = MulticlassClassificationEvaluator(
    labelCol="index_label", predictionCol="prediction", metricName="accuracy")
accuracy_percentage = multi_class_classification_evaluator.evaluate(predictions_test)

print("Test set error = %f " % (1.0 - accuracy_percentage))

```

Created random forest classifier with :

RandomForestClassificationModel (uid=RandomForestClassifier_5f69f240db35) with 30 trees

```

+-----+-----+-----+
|predicted_label|label|          features|
+-----+-----+-----+
|          5.0|  5.0|(22,[0,1,2,3,4,5,...|
|          5.0|  5.0|(22,[0,1,2,3,4,5,...|
|          5.0|  5.0|(22,[0,1,2,3,4,5,...|
|          5.0|  5.0|(22,[0,1,2,3,4,5,...|
|          5.0|  5.0|(22,[0,1,2,3,4,5,...|
|          5.0|  5.0|(22,[0,1,2,3,4,5,...|

```



```
|          5.0|  5.0|(22,[0,1,2,3,4,5,...|
|          5.0|  5.0|(22,[0,1,2,3,4,5,...|
|          5.0|  5.0|(22,[0,1,2,3,4,5,...|
|          5.0|  5.0|(22,[0,1,2,3,4,5,...|
+-----+-----+-----+-----+-----+
```

only showing top 10 rows

Test set error = 0.436635

CPU times: user 46 ms, sys: 11 ms, total: 57 ms

Wall time: 35.3 s

One-vs-Rest classifier (a.k.a. One-vs-All)

Create classifier with multiple binary classifiers and combine

Increasing number of iterations in this case doesn't affects accuracy beyond 39%

```
In [10]: %%time
#split train and test set in 80:20 ratio
(train, test) = data.randomSplit([0.8, 0.2])

# instantiate logistic regression classifier
logistic_regression_classifier = LogisticRegression(maxIter=30,
                                                    tol=1E-7, fitIntercept=True)

# initializer our one vs rest classifier
one_vs_rest = OneVsRest(classifier=logistic_regression_classifier)

# train model on training data
one_vs_rest_model = one_vs_rest.fit(train)

# get predictions on test data
test_predictions = one_vs_rest_model.transform(test)

# Test accuracy using multi class evaluator
multi_class_evaluator = MulticlassClassificationEvaluator(metricName="accuracy")
accuracy_percentage = multi_class_evaluator.evaluate(test_predictions)
print("Test set error = %f " % (1.0 - accuracy_percentage))
```

Test set error = 0.614291

CPU times: user 256 ms, sys: 58 ms, total: 314 ms

Wall time: 43.2 s

Feedforward Artificial Neural Network Classifier

- References :

- <https://dzone.com/articles/deep-learning-via-multilayer-perceptron-classifier>

- <https://spark.apache.org/docs/2.2.1/api/java/index.html?org/apache/spark/ml/classification/Random>

Neural network test test accuracy calculation has some issues with library

```
In [12]: %%time
#split data in 60:40 ratio
random_split = data.randomSplit([0.6, 0.4])
train = random_split[0]
test = random_split[1]

# specify input, middle, output layers
# No of input layers - 20
# intermediate layers - 15, 10, 8
# output layers - 3
layers = [20, 15, 10, 8, 6]

# create multilayer neural network classifier
nn_model = MultilayerPerceptronClassifier(maxIter=100,
                                           layers=layers, blockSize=128)

# train our neural network
model = nn_model.fit(train)

# compute accuracy on the test set
test_predictions = model.transform(test)

predictionAndLabels = test_predictions.select("prediction", "label")
```

CPU times: user 31.3 ms, sys: 6.53 ms, total: 37.9 ms

Wall time: 28.7 s

```
1 # Team Members
2 # - Shubhendu Vimal - 11915067
3 # - Dharani Kiran Kavuri - 11915033
4 # - Anmol More - 11915043
5
6 #Data Preparation in plain Python
7
8 #Read raw fasta protien sequence files and convert it to parsable CSV files
... using biopython
9
10 #Embedded block of code for converting to pdf only, run separately (time
... taking)
11
12 import sys
13 import pandas as pd
14 from Bio.SeqUtils.ProtParam import ProteinAnalysis
15
16 #arg 1 - fasta file name
17 #arg 2 - no of records to process
18 file = str(sys.argv[1])
19 limit = int(sys.argv[2])
20 print("Processing file : ", file)
21
22 df = pd.DataFrame()
23 enzyme_type = file[:-6]
24
25 # read fasta files, line by line and process enzyme sequence
26 print("Group all as : ", enzyme_type)
27 with open(file) as fileobject:
28     start = True
29     row = ""
30     for line in fileobject:
31         if("sp|" in line or "tr|" in line) :
32             start = True
33             if(row != "") :
34                 row_dict = {}
35                 row_dict["Sequence"] = row
36                 row_dict["Type"] = enzyme_type
37
38             try :
39                 #use biopython library to process enzyme sequence
40                 analysed_seq = ProteinAnalysis(row)
41                 #print(analysed_seq)
42                 amino_acid_counts = analysed_seq.count_amino_acids()
43                 row_dict.update(amino_acid_counts)
44                 analysed_seq = ProteinAnalysis(row)
45                 row_dict["weight"] = analysed_seq.molecular_weight()
46                 row_dict["gravy"] = analysed_seq.gravy()
47                 df = df.append(row_dict, ignore_index=True)
48                 print(df.shape)
49                 if(df.shape[0] > limit) :
50                     break
```

```
51         except :
52             print("Error")
53             row = ""
54         elif(start) :
55             row += line.rstrip()
56     print(df.head())
57 df.to_csv("data/" + enzyme_type + ".csv", index=False)
```